

- 13.4 Consider the relations  $r_1(A, B, C)$ ,  $r_2(C, D, E)$ , and  $r_3(E, F)$ , with primary keys  $A$ ,  $C$ , and  $E$ , respectively. Assume that  $r_1$  has 1000 tuples,  $r_2$  has 1500 tuples, and  $r_3$  has 750 tuples. Estimate the size of  $r_1 \bowtie r_2 \bowtie r_3$ , and give an efficient strategy for computing the join.

We could use the associative rule:

$$\begin{aligned} r_1 \bowtie r_2 \bowtie r_3 &= ((r_1 \bowtie r_2) \bowtie r_3) \quad ① \\ &= (r_1 \bowtie (r_2 \bowtie r_3)) \quad ② \end{aligned}$$

Since Primary Keys are  $A, C$ , and  $E$ , it could be concluded that

$$\text{tuples} = \min(\max(\min(1000, 1500), 750), \max(1000, \min(1500, 750)))$$

Therefore, we could use solution ①, the size is 1000.

Primary Key  $C$  could be index for  $r_2$ .

Primary Key  $E$  could be index for  $r_3$ .

For each tuple in  $r_1$ , we could use index  $C$  of  $r_2$  to match attribute  $C$  in  $r_1$ .

and we could use index  $E$  of  $r_3$  to match attribute  $E$  in  $r_2$ .

- 13.5 Consider the relations  $r_1(A, B, C)$ ,  $r_2(C, D, E)$ , and  $r_3(E, F)$  of Practice Exercise 13.4. Assume that there are no primary keys, except the entire schema. Let  $V(C, r_1)$  be 900,  $V(C, r_2)$  be 1100,  $V(E, r_2)$  be 50, and  $V(E, r_3)$  be 100. Assume that  $r_1$  has 1000 tuples,  $r_2$  has 1500 tuples, and  $r_3$  has 750 tuples. Estimate the size of  $r_1 \bowtie r_2 \bowtie r_3$  and give an efficient strategy for computing the join.

For each tuple of  $r_1$ .  $\frac{1500}{V(C, r_2)} = \frac{15}{11}$  tuples of  $r_2$  will join it.

For each tuple of  $r_2$ .  $\frac{1000}{V(C, r_1)} = \frac{10}{9}$  tuples of  $r_1$  will join it.

$\frac{750}{V(E, r_3)} = \frac{15}{2}$  tuples of  $r_3$  will join it.

For each tuple of  $r_3$ .  $\frac{1500}{V(E, r_2)} = 30$  tuples of  $r_2$  will join it.

Since  $r_1 \bowtie r_2 \bowtie r_3 = (r_1 \bowtie r_2) \bowtie r_3 \quad ①$

$$= r_1 \bowtie (r_2 \bowtie r_3) \quad ②$$

For solution ①  $\text{total}_1 = 1000 \times \frac{15}{11} \times \frac{15}{2} = 10228$

For solution ②  $\text{total}_2 = 750 \times 30 \times \frac{10}{9} = 25000$ ,  $10228 < 25000$ .

Therefore, join  $r_1$  and  $r_2$  first, then join  $r_3$ .

14.12 List the ACID properties. Explain the usefulness of each.

① Atomicity (原子性): Couldn't be divided.

Either all the operations are conducted properly or none are.

Usefulness: Keeping consistency.

② Consistency (一致性): Execution of a transaction is isolation preserves the consistency of database.

Usefulness: Programmers should write SQL correctly.

③ Isolation (孤立性): Each transaction is unaware of another transaction(s) while other transactions are acting concurrently.

Usefulness: Ensure that transaction won't be affected by wrong transactions.

④ Durability (可持续性): After a transaction completes, the changes of the database consist even if the operation is wrong or system failures.

14.15 Consider the following two transactions:

$T_{13}$ : read(A);  
read(B);  
if A = 0 then B := B + 1;  
write(B).  
 $T_{14}$ : read(B);  
read(A);  
if B = 0 then A := A + 1;  
write(A).

Table.

A	B	0VB
0	0	0 ✓
0	1	1 ✓
1	0	0 ✓
1	1	1 X

Let the consistency requirement be  $A = 0 \vee B = 0$ , with  $A = B = 0$  the initial values.

- Show that every serial execution involving these two transactions preserves the consistency of the database.
- Show a concurrent execution of  $T_{13}$  and  $T_{14}$  that produces a nonserializable schedule.
- Is there a concurrent execution of  $T_{13}$  and  $T_{14}$  that produces a serializable schedule?

a. ①  $T_{13}, T_{14}$

Initially A B 0VB

After  $T_{13}$  0 1 1

After  $T_{14}$  0 1 1

$A = 0 \vee B = 0 \equiv T \vee F = T$

②  $T_{14}, T_{13}$

Initially A B 0VB

After  $T_{14}$  1 0 0

After  $T_{13}$  1 0 0

$A = 0 \vee B = 0 \equiv F \vee T = F$

b.

$T_{13}$   
read(A) A=0

read(B) B=0  
if A=0: B:=B+1

write(B) B=1

$A = 0 \vee B = 0 \equiv F \vee F = F$

$T_{14}$ .

read(B) B=0

read(A) A=0

if B=0: A:=A+1

write(A) A=1

c. There exists no concurrent solution. Here is the reason:

Table.

A	B	0 V B	
0	0	0	✓
0	1	1	✓
1	0	0	✓
1	1	1	X

According to the table, only when  $A \neq B$  or  $A=B=0$  can be satisfied.

There exists only two ways:

① Start  $T_{14}$  after  $T_{13}$  is finished.

② Start  $T_{13}$  after  $T_{14}$  is finished.

The other circumstances, according to ACID. Isolation will cause  $A=0$  &  $B=0$  all the time. when  $T_{13}$  is reading or  $T_{14}$  is reading which will cause  $A=1$  &  $B=1$  in the end.

### Assignment

- A scheduled flight has 50 tickets. Agency one wants to book 10 tickets and agency two want to book 20 tickets. Please design a lock strategy to implement the concurrency control.

Agency one	Agency two
Read(A)	
	Read(A)
$A=A-10$	
Write(A)	
	$A=A-20$
	Write(A)

Agency One.

lock-S(A)  
lock-X(A)  
Read(A)  
 $A = A - 10$   
Write(A)  
unlock-X(A)  
unlock-S(A)

Agency Two.

lock-S(A)  
Read(A)  
 $A = A - 20$   
Write(A)  
unlock-S(A).

16.1 Explain why log records for transactions on the undo-list must be processed in reverse order, whereas redo is performed in a forward direction.

undo-list: undo is utilized to rollback the information. Backward direction could meet the requirement.

redo-list: redo is utilized to backup the information. Forward direction could meet the requirement.

If there exist an update of character 'a' → 'b' → 'c'

	backward	forward
undo	'c' → 'b' → 'a' ✓	'a' → 'b' ✗ initial value is 'a', false
redo	'c' → 'b' ✗ incorrect	'a' → 'b' → 'c' ✓

16.18 Consider the log in Figure 16.5. Suppose there is a crash just before the  $\langle T_0 \text{ abort} \rangle$  log record is written out. Explain what would happen during recovery.




Figure 16.5 Example of logged actions, and actions during recovery.

① Redo Pass (备份)

- Undo-List  $T_0, T_1$
- Start from  $\langle \text{checkpoint } \{T_0, T_1\} \rangle$
- Redo  $T_1$   $C=600$
- Undo-List  $T_0$
- Undo-List  $T_0, T_2$
- $A=400, B=2000$ .

② Undo Pass (回滚)

- Undo-List  $T_0, T_2$
- Roll back  $T_2$
- $A=500$ , output  $\langle T_2, A, 500 \rangle$
- output  $\langle T_2, \text{abort} \rangle$
- Roll back  $T_0$ .
- $B=2000$ , output  $\langle T_0, B, 2000 \rangle$
- output  $\langle T_0, \text{abort} \rangle$

Finally:

Result:  $A=400$   
 $B=2000$   
 $C=600$

Output:

$\langle T_2, A, 500 \rangle$   
 $\langle T_2, \text{abort} \rangle$   
 $\langle T_0, B, 2000 \rangle$   
 $\langle T_0, \text{abort} \rangle$